

BIN371_m3

2025-08-18

1) Objective

We evaluate predictive models on the **prepared train/test splits** for four datasets:

- Anthropometry
- ARI Symptoms
- Literacy
- Maternal Mortality

Target is **Value_scaled** (numeric) → **regression**.

We fit **Linear Regression (LM)**, **Decision Tree (rpart)**, and **Random Forest (RF)** on each dataset's **train** set, evaluate on the **test** set using **RMSE**, **MAE**, **R²**, and choose the **best model by lowest RMSE**. Diagnostics include residual plots and RF feature importance.

2) Inputs & Paths

Provide the absolute paths to your **Training** and **Test** folders (Windows paths may use forward slashes).

```
# --- where your train/test RDS live ---
train_dir <- "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Training-data"
test_dir  <- "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Test-data_m2"

# --- resolve the directory of this Rmd (works in Knit and in the editor) ---
rmd_path <- tryCatch(
  rstudioapi::getSourceEditorContext()$path,
  error = function(e) knitr::current_input()
)
rmd_dir <- if (nzchar(rmd_path)) normalizePath(dirname(rmd_path)) else normalizePath(getwd())

# --- output folder next to the Rmd ---
out_dir <- file.path(rmd_dir, "outputs_m3")
if (!dir.exists(out_dir)) dir.create(out_dir, recursive = TRUE)

# --- dataset names and RDS paths ---
datasets <- c("Anthropometry", "ARI_Symptoms", "Literacy", "MaternalMortality")

train_paths <- setNames(file.path(train_dir, paste0("train_", datasets, ".rds")), datasets)
test_paths  <- setNames(file.path(test_dir,  paste0("test_",  datasets, ".rds")),  datasets)

# sanity check
out_dir
```

```
## [1] "C:\\Users\\Suhil Jugroop\\OneDrive\\Documents\\GitHub\\BIN371-Project\\Milestone 3/outputs_m3"
```

```
train_paths
```

```
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Training-data_m2/tr  
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Training-data_m2/tr  
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Training-data_m  
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Training-data_m2/train_1
```

```
test_paths
```

```
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Test-data_m2/test_A  
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Test-data_m2/test_1  
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Test-data_m2/t  
##  
## "C:/Users/Suhil Jugroop/OneDrive/Documents/GitHub/BIN371-Project/Milestone 3/Test-data_m2/test_Matern
```

3) Helper Functions (Metrics, Fitters, Predictors)

```
# --- metrics ---  
rmse <- function(truth, pred) sqrt(mean((truth - pred)^2, na.rm = TRUE))  
mae <- function(truth, pred) mean(abs(truth - pred), na.rm = TRUE)  
r2 <- function(truth, pred) {  
  ss_res <- sum((truth - pred)^2, na.rm = TRUE)  
  ss_tot <- sum((truth - mean(truth, na.rm = TRUE))^2, na.rm = TRUE)  
  1 - ss_res/ss_tot  
}  
  
# --- coerce train/test list(X, y) to data.frame with Value_scaled ---  
as_df_xy <- function(obj) {  
  stopifnot(is.list(obj), all(c("X", "y") %in% names(obj)))  
  df <- as.data.frame(obj$X)  
  df$Value_scaled <- as.numeric(obj$y)  
  df  
}  
  
# --- model fitting for regression ---  
fit_models <- function(df_train) {  
  # defensive: drop columns with Inf/NaN  
  bad <- vapply(df_train, function(v) any(is.infinite(v) | is.nan(v)), logical(1))  
  if (any(bad)) df_train <- df_train[, !bad, drop = FALSE]  
  
  form <- as.formula("Value_scaled ~ .")
```

```

mod_lm <- tryCatch(lm(form, data = df_train), error = function(e) NULL)

ctrl <- rpart.control(minsplit = 10, cp = 0.001, maxdepth = 10, xval = 5)
mod_dt <- tryCatch(rpart(form, data = df_train, method = "anova", control = ctrl), error = function(e) NULL)

mod_rf <- tryCatch(randomForest(form, data = df_train, ntree = 500), error = function(e) NULL)

list(lm = mod_lm, dt = mod_dt, rf = mod_rf)
}

safe_predict <- function(model, newdata) {
  if (is.null(model)) return(rep(NA_real_, nrow(newdata)))
  as.numeric(tryCatch(predict(model, newdata = newdata), error = function(e) rep(NA_real_, nrow(newdata))))
}

evaluate_models <- function(models, df_test) {
  truth <- df_test$Value_scaled
  tibble(
    Model = c("Linear Regression", "Decision Tree", "Random Forest"),
    RMSE = c(
      rmse(truth, safe_predict(models$lm, df_test)),
      rmse(truth, safe_predict(models$dt, df_test)),
      rmse(truth, safe_predict(models$rf, df_test))
    ),
    MAE = c(
      mae(truth, safe_predict(models$lm, df_test)),
      mae(truth, safe_predict(models$dt, df_test)),
      mae(truth, safe_predict(models$rf, df_test))
    ),
    R2 = c(
      r2(truth, safe_predict(models$lm, df_test)),
      r2(truth, safe_predict(models$dt, df_test)),
      r2(truth, safe_predict(models$rf, df_test))
    )
  ) %>% arrange(RMSE)
}

```

4) Load Train/Test Splits

note this will produce a ton of empty data, don't worry about it, that's just how these things go when converted to binary, there is a way to hide it in the code apparently but I'm lazy

```

train_objs <- lapply(train_paths, readRDS)
test_objs <- lapply(test_paths, readRDS)

train_df <- lapply(train_objs, as_df_xy)
test_df <- lapply(test_objs, as_df_xy)

# quick preview
purrr::imap(train_df, ~{cat("\nTRAIN:", .y, "\n"); print(glimpse(.x))})
purrr::imap(test_df, ~{cat("\nTEST : ", .y, "\n"); print(glimpse(.x))})

```

5) Train Models Per Dataset

```
models <- lapply(train_df, fit_models)

# status report
purrr::imap(models, ~{
  cat("\n", .y, " models:\n", sep = "")
  cat("  LM:", if (is.null(.x$lm)) "NULL" else "OK", "\n")
  cat("  DT:", if (is.null(.x$dt)) "NULL" else "OK", "\n")
  cat("  RF:", if (is.null(.x$rf)) "NULL" else "OK", "\n")
})
```

```
##
## Anthropometry models:
##   LM: OK
##   DT: OK
##   RF: NULL
##
## ARI_Symptoms models:
##   LM: OK
##   DT: OK
##   RF: NULL
##
## Literacy models:
##   LM: OK
##   DT: OK
##   RF: NULL
##
## MaternalMortality models:
##   LM: OK
##   DT: OK
##   RF: NULL

## $Anthropometry
## NULL
##
## $ARI_Symptoms
## NULL
##
## $Literacy
## NULL
##
## $MaternalMortality
## NULL
```

6) Evaluate & Choose Best Model (by RMSE)

```
results <- purrr::imap_dfr(models, ~{
  ds <- .y
  ev <- evaluate_models(.x, test_df[[ds]])
```

```

ev$Dataset <- ds
ev
})

results_ranked <- results %>%
  group_by(Dataset) %>%
  arrange(RMSE, .by_group = TRUE) %>%
  mutate(Rank = row_number()) %>%
  ungroup()

knitr::kable(results_ranked, digits = 4, caption = "Model Comparison per Dataset (sorted by RMSE)")

```

Table 1: Model Comparison per Dataset (sorted by RMSE)

Model	RMSE	MAE	R2	Dataset	Rank
Linear Regression	0.3856	0.1906	0.8006	ARI_Symptoms	1
Decision Tree	0.8916	0.7969	-0.0659	ARI_Symptoms	2
Random Forest	NaN	NaN	1.0000	ARI_Symptoms	3
Decision Tree	0.8612	0.7427	-0.0470	Anthropometry	1
Linear Regression	0.9238	0.3872	-0.2048	Anthropometry	2
Random Forest	NaN	NaN	1.0000	Anthropometry	3
Decision Tree	0.4938	0.4903	-0.1139	Literacy	1
Linear Regression	0.5561	0.3072	-0.4123	Literacy	2
Random Forest	NaN	NaN	1.0000	Literacy	3
Decision Tree	0.7319	0.6432	-0.0030	MaternalMortality	1
Linear Regression	1.4526	0.8965	-2.9506	MaternalMortality	2
Random Forest	NaN	NaN	1.0000	MaternalMortality	3

```

best_models <- results_ranked %>%
  filter(Rank == 1) %>%
  select(Dataset, Best_Model = Model, RMSE, MAE, R2)

knitr::kable(best_models, digits = 4, caption = "Chosen Model per Dataset (lowest RMSE)")

```

Table 2: Chosen Model per Dataset (lowest RMSE)

Dataset	Best_Model	RMSE	MAE	R2
ARI_Symptoms	Linear Regression	0.3856	0.1906	0.8006
Anthropometry	Decision Tree	0.8612	0.7427	-0.0470
Literacy	Decision Tree	0.4938	0.4903	-0.1139
MaternalMortality	Decision Tree	0.7319	0.6432	-0.0030

7) Diagnostics (Residuals, RF Importance, Tree Plot)

```

plot_residuals <- function(model, df_test, title){
  pred <- safe_predict(model, df_test)
  res <- df_test$Value_scaled - pred
  p1 <- ggplot(data.frame(Pred = pred, Resid = res), aes(Pred, Resid)) +

```

```

    geom_point(alpha = 0.6) + geom_hline(yintercept = 0, linetype = "dashed") +
    labs(title = paste0(title, " - Residuals vs Predicted"), x = "Predicted", y = "Residuals")
  p2 <- ggplot(data.frame(Resid = res), aes(Resid)) +
    geom_histogram(bins = 30, color = "black") +
    labs(title = paste0(title, " - Residual Histogram"), x = "Residual", y = "Count")
  list(p1 = p1, p2 = p2)
}

plot_rf_importance <- function(rf_model, title, top_n = 15){
  if (is.null(rf_model) || is.null(rf_model$importance)) return(NULL)
  imp <- as.data.frame(rf_model$importance)
  imp$Feature <- rownames(imp)
  imp <- imp %>% arrange(desc(IncNodePurity)) %>% head(top_n)
  ggplot(imp, aes(x = reorder(Feature, IncNodePurity), y = IncNodePurity)) +
    geom_col() + coord_flip() +
    labs(title = paste0(title, " - RF Variable Importance (Top ", top_n, ")"), x = "Feature", y = "IncN
}

purrr::pwalk(list(best_models$Dataset, best_models$Best_Model), function(ds, mlabel){
  cat("\n## Diagnostics:", ds, "(", mlabel, ")\n")
  mdl <- switch(mlabel,
    "Linear Regression" = models[[ds]]$lm,
    "Decision Tree"     = models[[ds]]$dt,
    "Random Forest"     = models[[ds]]$rf)
  testd <- test_df[[ds]]

  # Residuals
  rp <- plot_residuals(mdl, testd, paste0(ds, " - ", mlabel))
  print(rp$p1); print(rp$p2)

  # Tree plot if DT exists
  if (!is.null(models[[ds]]$dt)) {
    rpart.plot(models[[ds]]$dt, main = paste0(ds, " - Decision Tree"))
  }

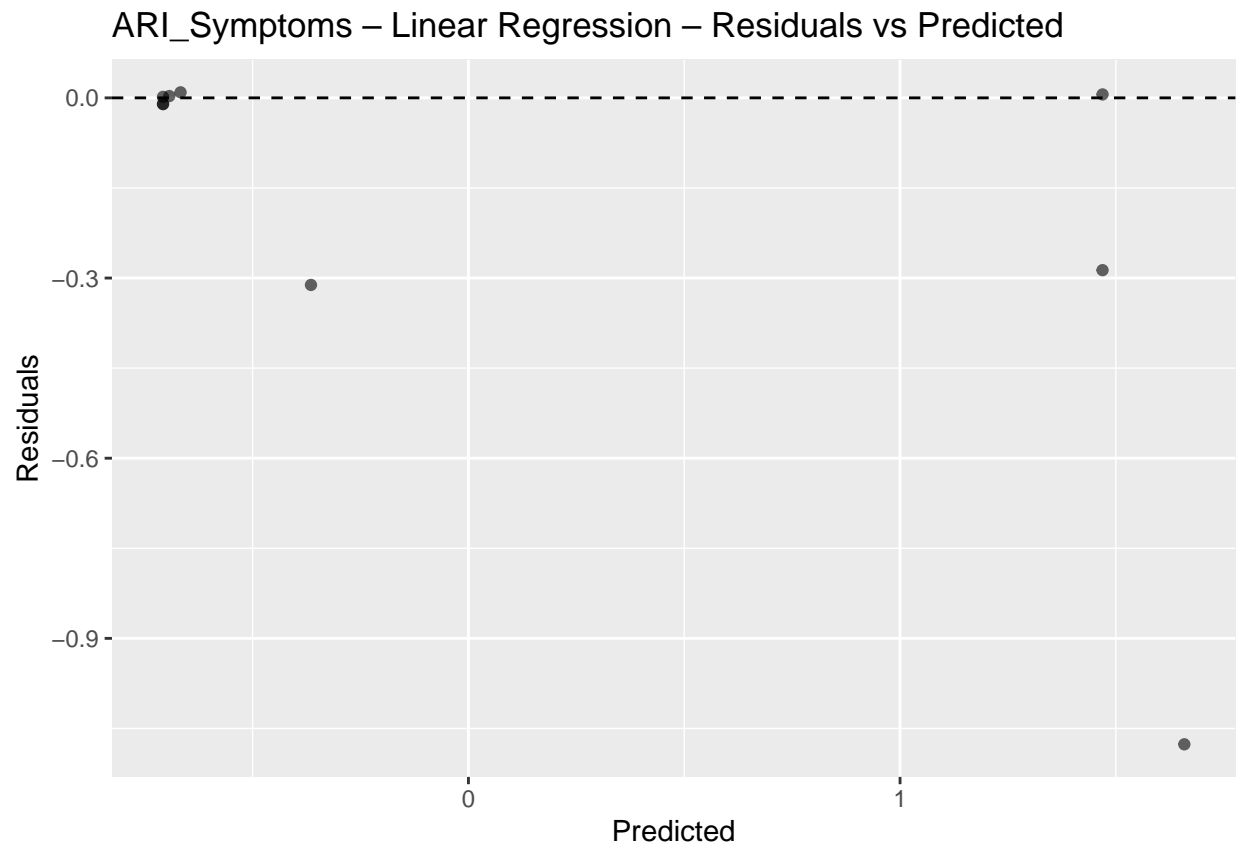
  # RF importance (always useful to show)
  if (!is.null(models[[ds]]$rf)) {
    vip <- plot_rf_importance(models[[ds]]$rf, ds, top_n = 15)
    if (!is.null(vip)) print(vip)
  }
})

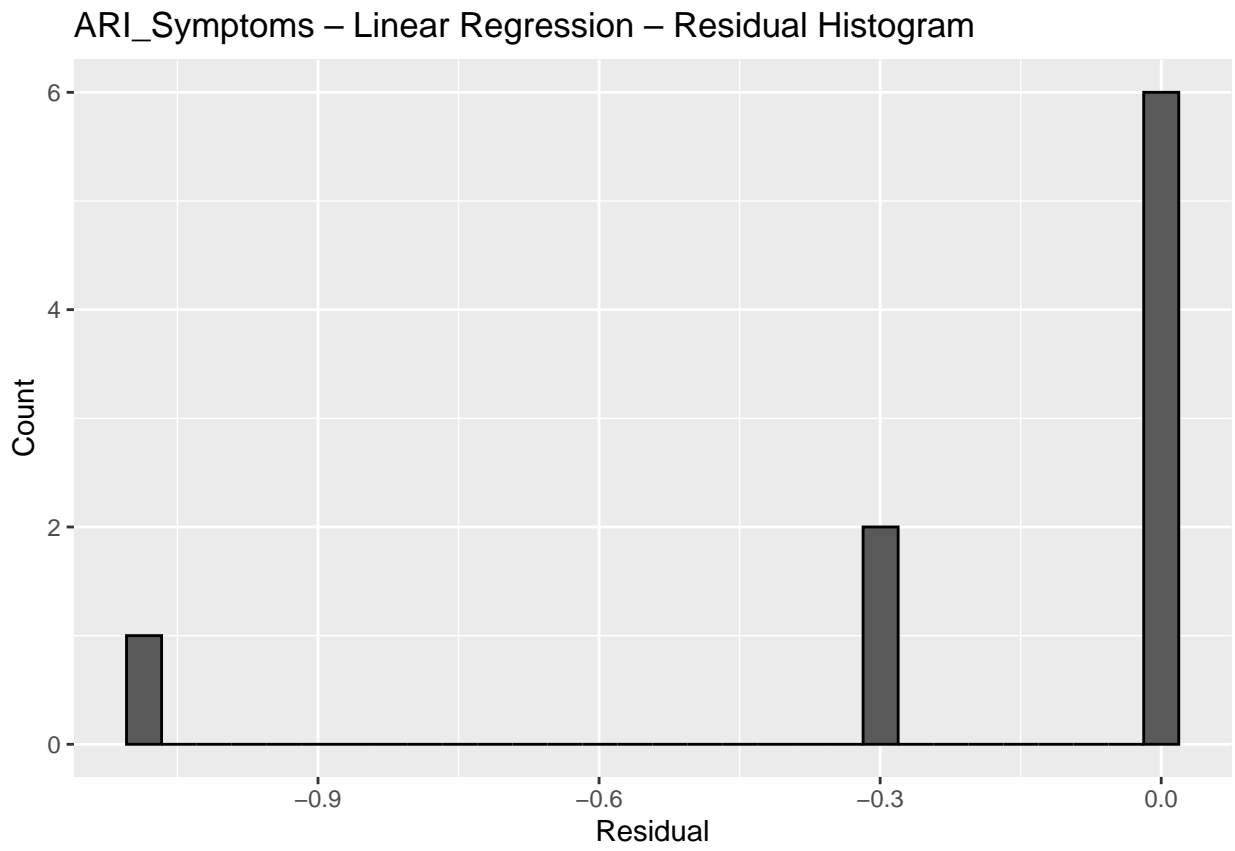
```

```

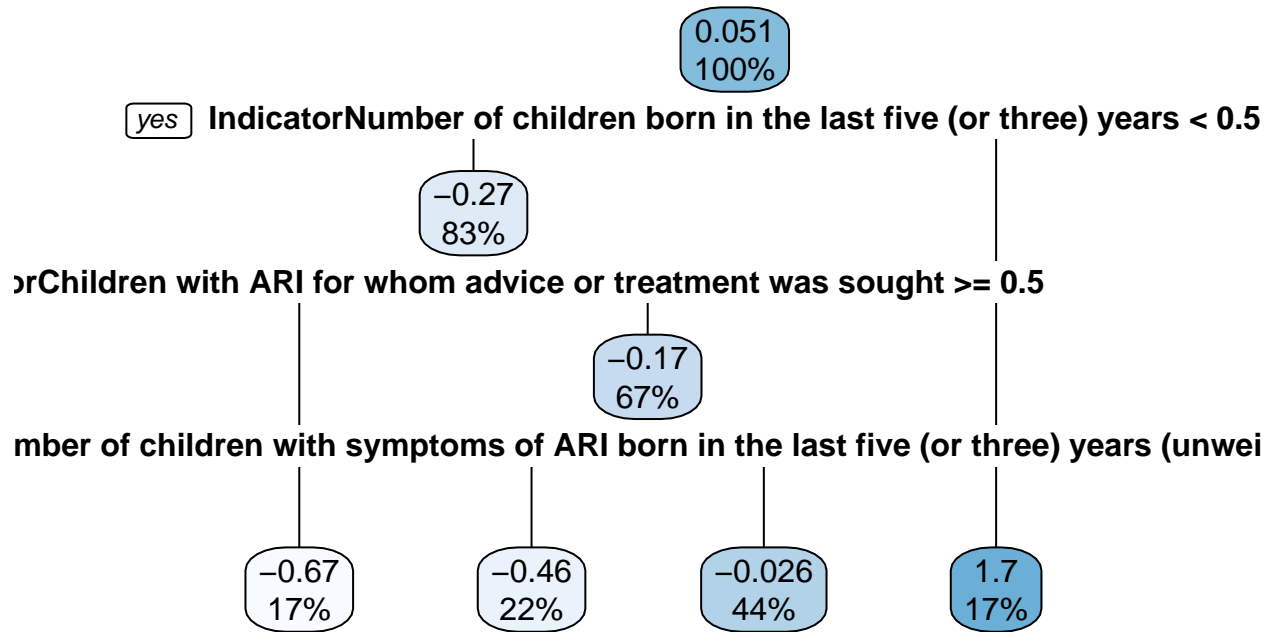
##
## ## Diagnostics: ARI_Symptoms ( Linear Regression )

```



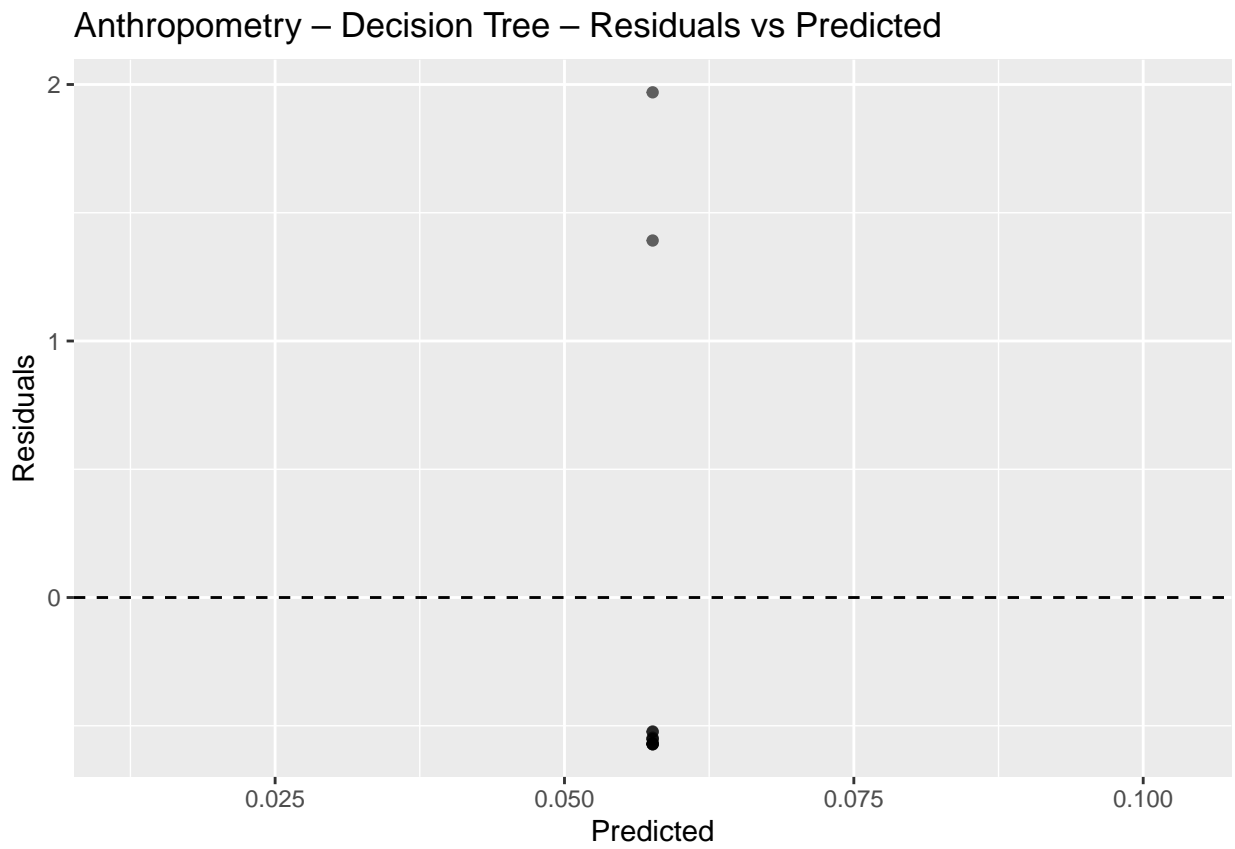


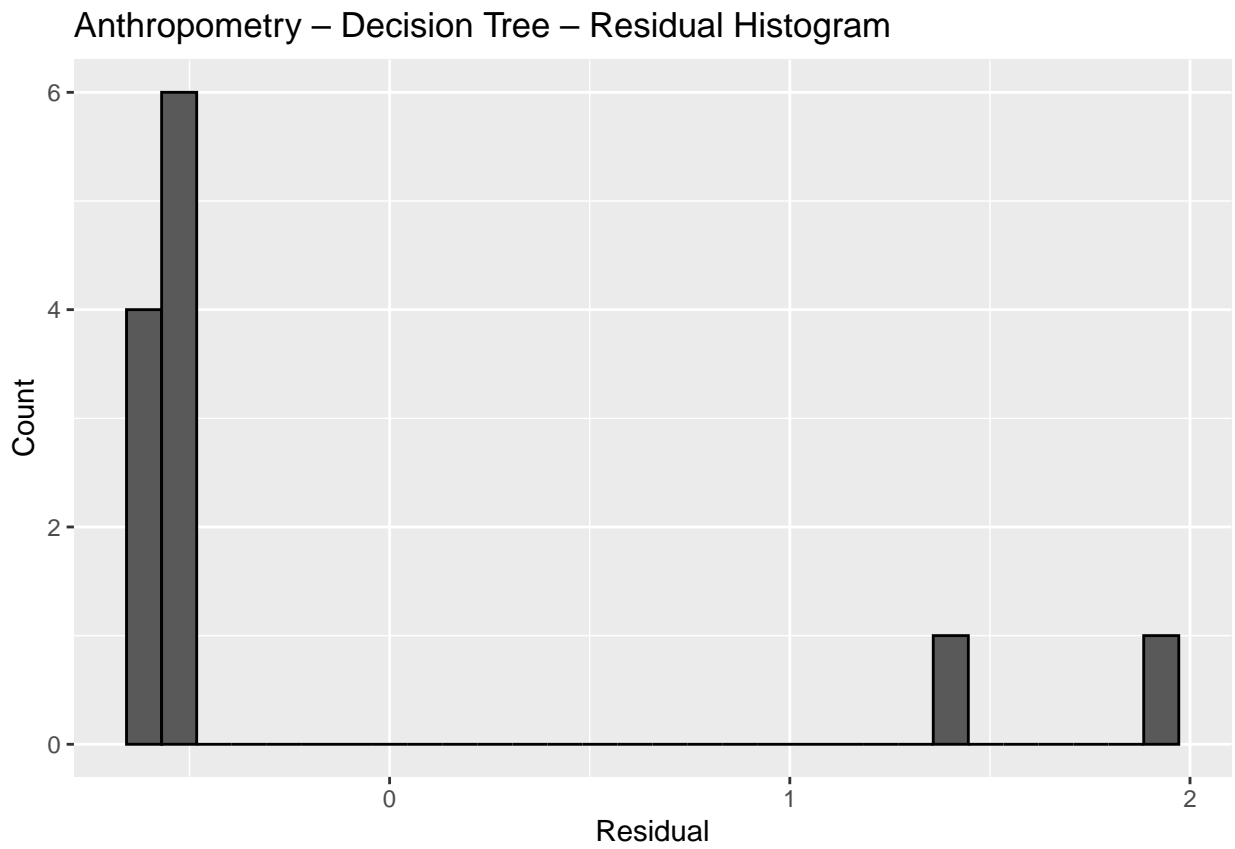
ARI_Symptoms – Decision Tree



##

Diagnostics: Anthropometry (Decision Tree)






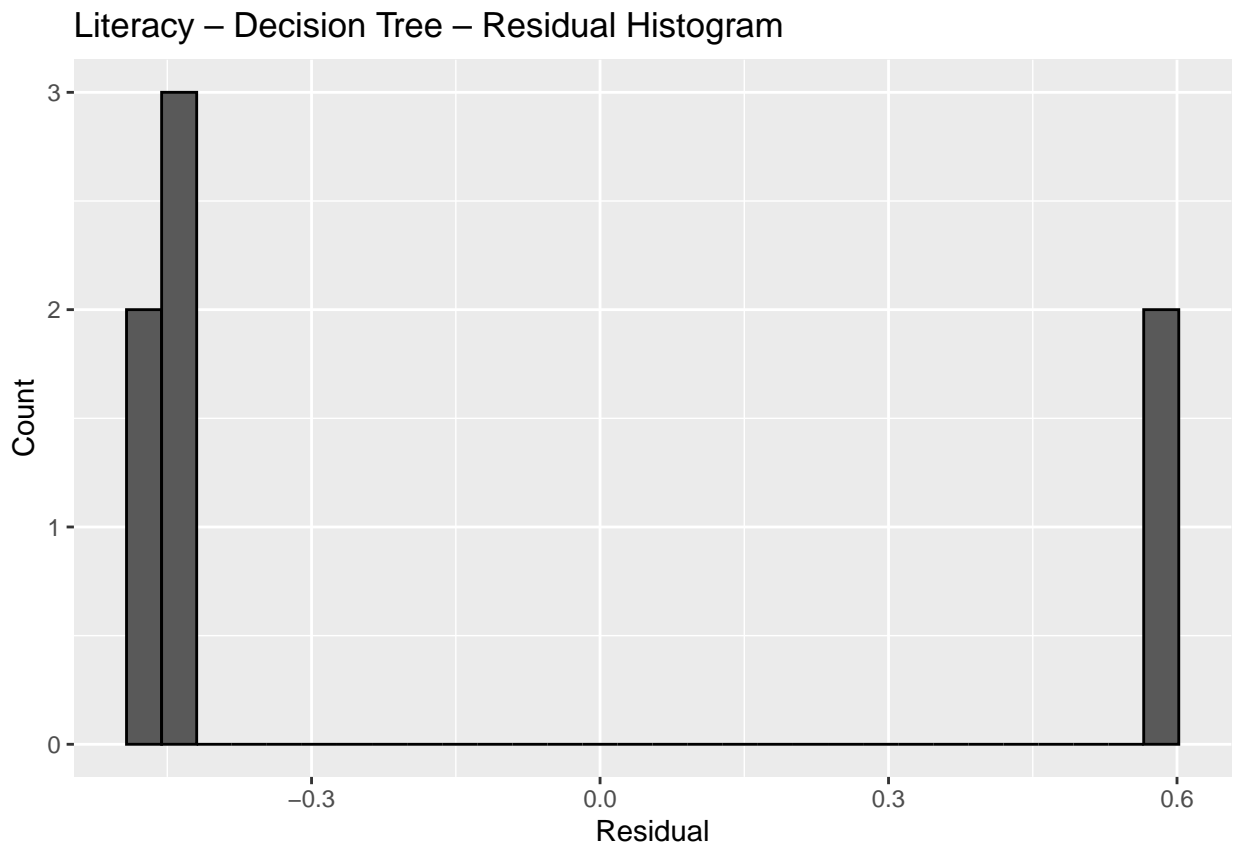
Anthropometry – Decision Tree

0.058
100%

##

Diagnostics: Literacy (Decision Tree)






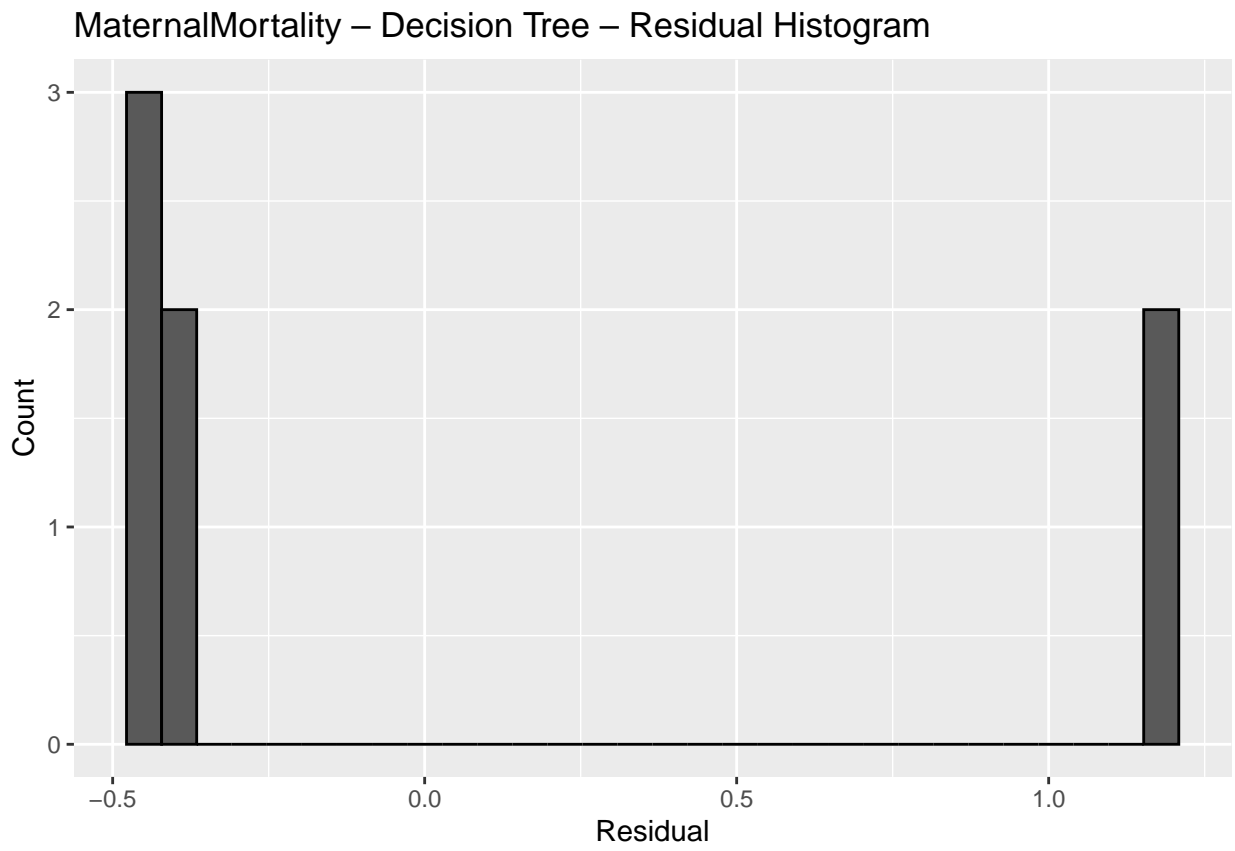
Literacy – Decision Tree

0.053
100%

##

Diagnostics: MaternalMortality (Decision Tree)





MaternalMortality – Decision Tree

-0.013
100%

8) Save Winning Models

```
save_winner <- function(ds, model_label) {  
  mdl <- switch(model_label,  
    "Linear Regression" = models[[ds]]$lm,  
    "Decision Tree"     = models[[ds]]$dt,  
    "Random Forest"     = models[[ds]]$rf)  
  if (!is.null(mdl)) {  
    fn <- file.path(out_dir, paste0("best_model_", gsub("[^A-Za-z0-9_]", "", ds), "_", gsub(" ", "", model_label), ".rds"))  
    saveRDS(mdl, fn)  
    cat("Saved:", fn, "\n")  
  }  
}  
purrr::pwalk(list(best_models$Dataset, best_models$Best_Model), save_winner)
```

```
## Saved: C:\Users\Suhil Jugroop\OneDrive\Documents\GitHub\BIN371-Project\Milestone 3/outputs_m3/best_model_Linear_Regression.rds  
## Saved: C:\Users\Suhil Jugroop\OneDrive\Documents\GitHub\BIN371-Project\Milestone 3/outputs_m3/best_model_Decision_Tree.rds  
## Saved: C:\Users\Suhil Jugroop\OneDrive\Documents\GitHub\BIN371-Project\Milestone 3/outputs_m3/best_model_Random_Forest.rds  
## Saved: C:\Users\Suhil Jugroop\OneDrive\Documents\GitHub\BIN371-Project\Milestone 3/outputs_m3/best_model_Linear_Regression.rds
```

9) Rubric Coverage (place in report)

- **Modelling Technique (10%)**: Regression task; techniques compared: LM, DT, RF; assumptions noted (LM linearity/homoscedasticity; RF nonparametric).
- **Test Design (10%)**: Predefined train/test splits; evaluation with RMSE/MAE/R²; fixed random seed.
- **Model Building (10%)**: Training code and parameters (`rpart.control`, `ntree=500` for RF).
- **Assess the Model (10%)**: Comparison table + best model table; diagnostics (residuals, tree, RF importance).
- **Visual Storytelling (20%)**: Use the generated plots; optionally port to Power BI/Shiny.
- **Reporting & Documentation (10%)**: Knit this Rmd to HTML/PDF and narrate results.
- **Oral Presentation (20%)**: Summarize best model per dataset, key features, recommendations.
- **Project Files & Code (10%)**: Include this Rmd, saved models in `outputs_m3/`, and figures.

```
cat("\nMilestone 3 complete.\nBest models saved to: ", normalizePath(out_dir), "\n")
```

```
##
```

```
## Milestone 3 complete.
```

```
## Best models saved to: C:\Users\Suhil Jugroop\OneDrive\Documents\GitHub\BIN371-Project\Milestone 3\o
```